

# DEEP LEARNING

Trainer : **Dr. Darshan Ingle**



# In-Depth Loss Functions



# 1. Mean Squared Error

- from a probabilistic perspective.
- helps prepare us for cross-entropy loss

Note: Error = Cost = Loss = Objective

$$MSE = \frac{1}{N} \cdot \sum_{i=1}^n (y_i - \hat{y}_i)^2$$


$$ME = \frac{-2 + 2}{2} = 0$$



# MSE- Why is it squared?




$[-4, +4]$



$$\text{Mean error} = \frac{(+4) + (+4) + (-4) + (-4)}{4} = \frac{0}{4} = 0$$

$$\text{Mean Absolute error} = \frac{|+4| + |+4| + |-4| + |-4|}{4} = \frac{4+4+4+4}{4} = 4$$



$$\text{MAE} = \frac{|+7| + |+1| + |-2| + |-6|}{4} = \frac{7+1+2+6}{4} = \frac{16}{4} = 4$$

③

$$\text{Mean Squared error} = \frac{(+4)^2 + (+4)^2 + (-4)^2 + (-4)^2}{4} = \frac{64}{4} = 16$$

$$= \frac{(+7)^2 + (+1)^2 + (-2)^2 + (-6)^2}{4} = \frac{49+1+4+36}{4} = \frac{90}{4} = 22.5$$

# MSE- Footnote 1

MSE  $\rightarrow$  It is very difficult to map to original values.

$\therefore$  We go for Root Mean Square error.

$$\boxed{RMSE = \sqrt{MSE}}$$

$$RMSE_1 = \sqrt{16} = 4$$

$$[-4, +4]$$

$$RMSE_2 = \sqrt{22.5} \approx 4.74$$

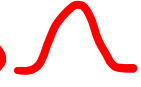
$$[-6, +7]$$

# MSE- Footnote 2

## 2. Maximum Likelihood Estimation (MLE)

- <https://www.youtube.com/watch?v=XepXtl9YKwc>

# MLE Formula

eg: We model the heights of the students in our class as a Gaussian distribution (Bell curve) 

$$\text{PDF} = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp\left(-\frac{1}{2} \cdot \frac{(x-\mu)^2}{\sigma^2}\right)$$



# MLE vs MSE Conclusion

Maximizing the likelihood is same as minimizing the squared error.  
i.e.  $(x_i - \mu)$

$$L = -\frac{1}{N} \cdot \sum_{i=1}^N (x_i - \mu)^2$$


$$MSE = +\frac{1}{N} \cdot \sum_{i=1}^N (x_i - \mu)^2$$

### 3. Cross-Entropy / Log Loss

- If you are training a binary classifier, chances are you are using binary cross-entropy / log loss as your loss function.
- Have you ever thought about what exactly does it mean to use this loss function?
- The thing is, given the ease of use of today's libraries and frameworks, it is very easy to overlook the true meaning of the loss function used.

# Cross-Entropy / Log Loss

# Cross-Entropy / Log Loss [0 - 1]



- Cross-entropy and log loss are slightly different depending on context, but in machine learning when calculating error rates between 0 and 1 they resolve to the same thing.

# Cross-Entropy / Log Loss

## Math

- In binary classification, where the number of classes  $M$  equals 2, cross-entropy can be calculated as:  
$$-(y \log(p) + (1-y) \log(1-p))$$
- If  $M > 2$  (i.e. multiclass classification), we calculate a separate loss for each class label per observation and sum the result.

$$\left\{ - \sum_{c=1}^M y_{o,c} \log(p_{o,c}) \right\}$$

- **Note:**
- $M$  - number of classes (dog, cat, fish)
- $\log$  - the natural log
- $y$  - binary indicator (0 or 1) if class label  $c$  is the correct classification for observation  $o$
- $p$  - predicted probability observation  $o$  is of class  $c$

# Cross-Entropy / Log Loss

Code:

```
def CrossEntropy(yHat, y):  
    if y == 1:  
        return -log(yHat)  
    else:  
        return -log(1 - yHat)
```


# Cross-Entropy / Log Loss

## A Simple Classification Problem


- Let's start with 10 random points:

$x = [-2.2, -1.4, -0.8, 0.2, 0.4, 0.8, 1.2, 2.2, 2.9, 4.6]$

This is our only feature:  $x$ .




Now, let's assign some colors to our points: red and green. These are our **labels**.



Green = 1.0  
Red = 0.0

# Cross-Entropy / Log Loss




So, our classification problem is quite straightforward: given our **feature  $x$** , we need to predict its **label: red or green**.

Since this is a **binary classification**, we can also pose this problem as: “**is the point green**” or, even better, “**what is the probability of the point being green**”? Ideally, **green points** would have a probability of **1.0** (of being green), while **red points** would have a probability of **0.0** (of being green).



# Cross-Entropy / Log Loss



If we fit a model to perform this classification, it will predict a probability of being green to each one of our points. Given what we know about the color of the points, how can we evaluate how good (or bad) are the predicted probabilities? This is the whole purpose of the loss function! It should return high values for bad predictions and low values for good predictions.

For a binary classification like our example, the typical loss function is the binary cross-entropy / log loss.

# Cross-Entropy / Log Loss

If you look this **loss function** up, this is what you'll find:

$$\left[ H_p(q) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i)) \right]$$

$$- \sum p \cdot \log p$$

where y is the label (1 for green points and 0 for red points) and p(y) is the predicted probability of the point being green for all N points.


Green, y=1

Red, y=0

# Cross-Entropy / Log Loss


- Computing the Loss — the visual way

*Split the points acc. to classes (pos or neg).*




# Cross-Entropy / Log Loss


Train a Log-Reg.



# Cross-Entropy / Log Loss




# Cross-Entropy / Log Loss




# Cross-Entropy / Log Loss

*put it altogether*



# Cross-Entropy / Log Loss

*We only & only need probabilities*





# Cross-Entropy / Log Loss




# Cross-Entropy / Log Loss


- Since we're trying to compute a loss, we need to penalize bad predictions, right? If the **probability** associated with the **true class** is **1.0**, we need its **loss** to be **zero**. Conversely, if that **probability is low**, say, **0.01**, we need its **loss** to be **HUGE!**
- It turns out, taking the (negative) log of the probability suits us well enough for this purpose *(since the log of values between 0.0 and 1.0 is negative, we take the negative log to obtain a positive value for the loss).*

# Cross-Entropy / Log Loss


- The plot below gives us a clear picture —as the predicted probability of the true class gets closer to zero, the loss increases exponentially:



# Cross-Entropy / Log Loss



Take -ve log of prob.



Binary Cross  
entropy / log loss

