


# DEEP LEARNING – Convolutional Neural Network

Trainer: Dr. Darshan Ingle.

# Data Augmentation




# Why data is important?

891 rows X 12 cols.

X ~ y 1/0





891  
+100 (Bias)



# Why data is important?

With images, we can invent new data & still its meaning will remain the same.

Also, no biases will be included.



# Problem with this approach

## Data takes up space




The more data I invent, the more space it takes up!

There are an endless number of ways I can invent new data

Shift to the left by 1, 2, 3, 4, 5, ... pixels

Right / up / down also

Rotation





# This can be done all automatically!

Have you ever considered how to rotate an image?

Use Tensorflow's Keras API instead

We need to know about generators / iterators



# Generators / Iterators

Loop from 0...10: for i in range(10)

0, 1, ..., 9

In Python 2, `range(10)` = `[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]` # a list!

In Python 2, use `xrange(10)` # does NOT create a list

In Python 3, `print(range(10))` yields `range(0, 10)`

- Not a list!


0  
1  
2  
3  
4  
5  
...

# Create your own generator

- Can you write your own function to do something like this?

```
for x in my_random_generator():  
    print(x)
```

```
# 0.06654313  
# -0.68315371  
# 1.46795401  
# -0.9017639  
# 0.77572637
```





# Yield

```
def my_random_generator():  
    for _ in range(10):  
        x = np.random.randn()  
        yield x
```

Notice: no list is ever created

All values do not need to be stored in memory simultaneously

# Apply this to Data Augmentation

You could similarly generate augmented data on the fly  
Conceptually, think of it like this

```
def my_image_augmentation_generator():  
    for x_batch, y_batch in zip(x_train, y_train):  
        x_batch = augment(x_batch)  
        yield x_batch, y_batch
```

~~return~~

1000 images

Batch Size = 100

$$\text{No. of Batches} = \frac{1000}{100} = 10$$

✓ B1 (x, y)  
B2 (x, y)  
B3 (x, y)  
⋮ ⋮  
B10 (x, y)

# How does it work in tf.keras?

① `from tensorflow.keras.preprocessing.image import`  
`ImageDataGenerator`

② `data_generator = ImageDataGenerator(  
width_shift_range=0.1,  
height_shift_range=0.1,  
horizontal_flip=True)`

Other args: rotation\_range, width\_shift\_range,  
height\_shift\_range, brightness\_range, shear\_range,  
zoom\_range, horizontal\_flip, and vertical\_flip

# How does it work in tf.keras?

```
data_generator = ImageDataGenerator(...)  
train_generator = data_generator.flow(  
    x_train, y_train, batch_size)
```


```
steps_per_epoch = x_train.shape[0] // batch_size  
r = model.fit_generator(  
    train_generator,  
    steps_per_epoch=steps_per_epoch,  
    epochs=50)
```

# Business as usual

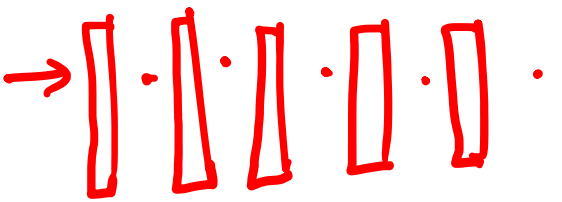
`fit_generator` returns history, so plot loss per iteration, etc.

```
# Plot loss per iteration
import matplotlib.pyplot as plt
plt.plot(r.history['loss'], label='loss')
plt.plot(r.history['val_loss'], label='val_loss')
plt.legend()
```


<matplotlib.legend.Legend at 0x7f92f7ca6908>



# Batch Normalization

Img → Norm. → 

- Early on, we noted that it's important to normalize/standardize data before passing it into algorithms like linear/logistic regression
- Problem: because this operation is done only on input data, only the first layer sees normalized data (after being transformed by Dense layer it's no longer normalized)



$$z = (x - \mu) / \sigma$$

How can we mk  
the data at every  
layer normalized?  
Soln: Batch  
Normalization


# Batch Normalization

- To start, recall: In Tensorflow, when we call model.fit(), we are doing batch gradient descent

```
for epoch in range(epochs):  
    for x_batch, y_batch in next_batch(x_train, y_train):  
        w ← w - learning_rate * grad(x_batch, y_batch)
```

# Batch Normalization

- What if we had a *layer* that would look at each batch, calculate the mean and standard deviation on the fly, and standardize based on that?



$$z = (x - \mu) / \sigma$$



# Batch Norm as Regularization


Can help with overfitting

Since every batch is slightly different, you'll get a slightly different  $\mu_B, \sigma_B$


They are not the true mean / std of the whole dataset

This is essentially noise, and using noise during training makes the neural network impervious to noise

- Rather than fitting to the noise!  
(a.k.a. "overfitting")



# Where is Batch Norm used?



Check out architecture of popular NN : VGG, ResNet, Inception, etc.

Refer NB “7  
CNN\_CIFAR\_Improved.ipynb”



# CNN Completed!!!