

Transfer learning with TF-hub

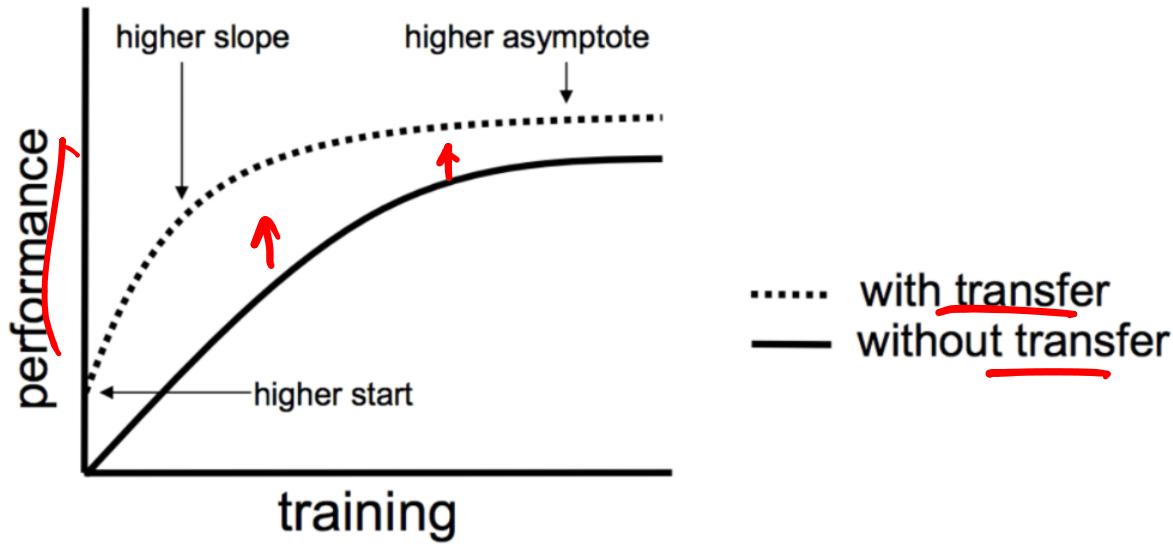
Instructor: Dr. Darshan
Ingle



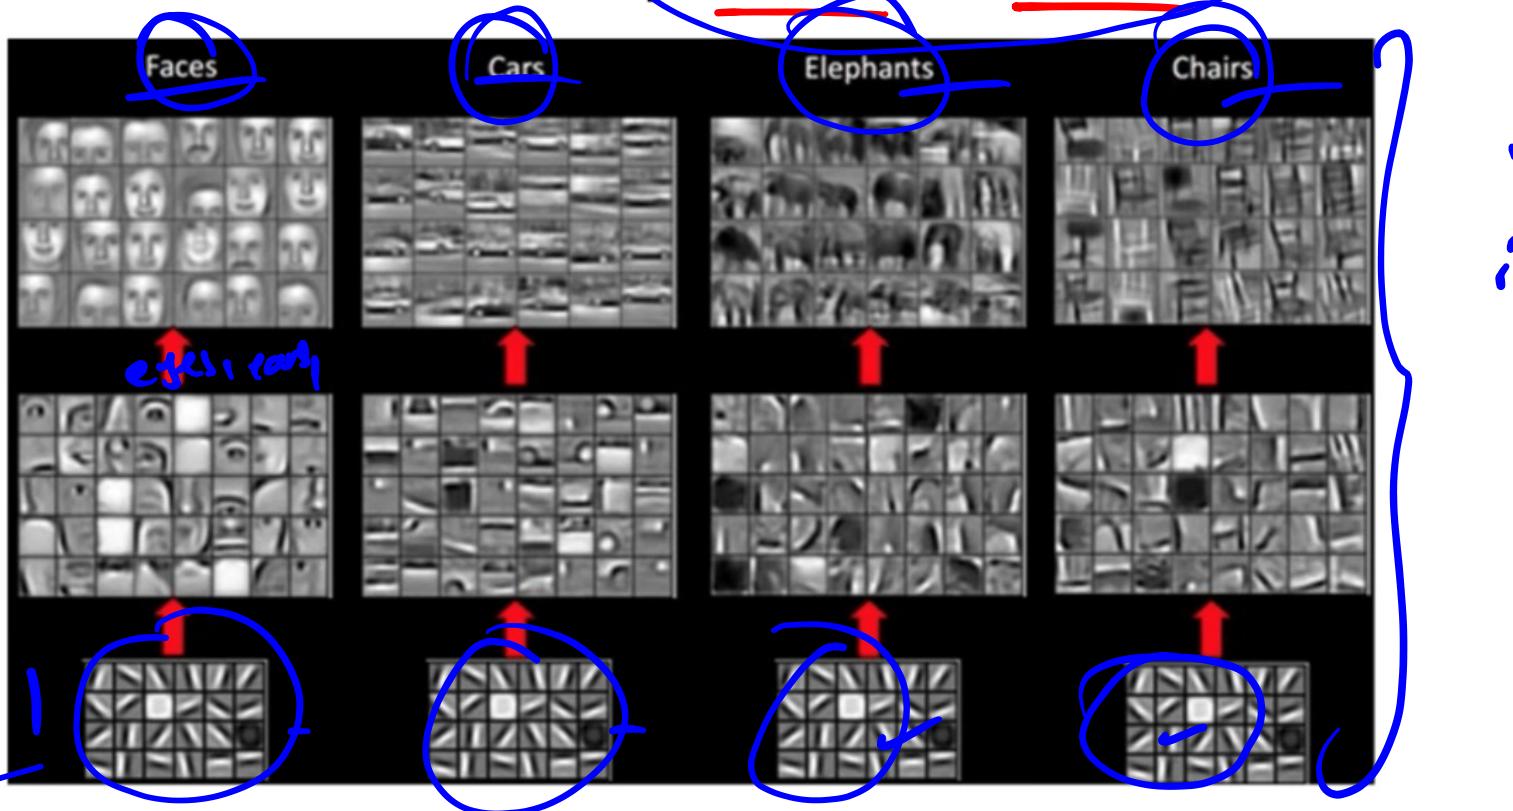
HI Y'ALL !!!

STARTING SOON

- A very important topic in modern deep learning



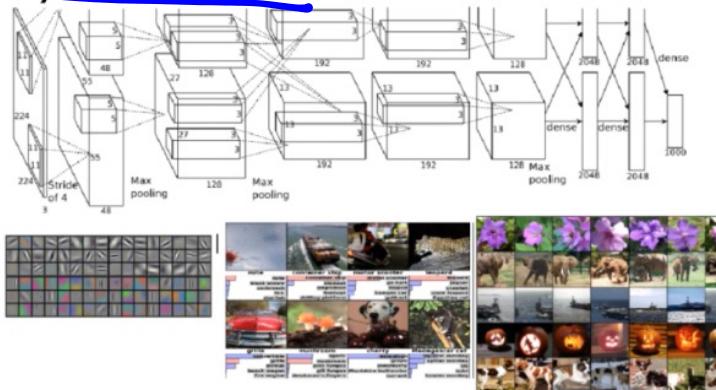
- Notice: *all* CNNs learn simple lines and strokes

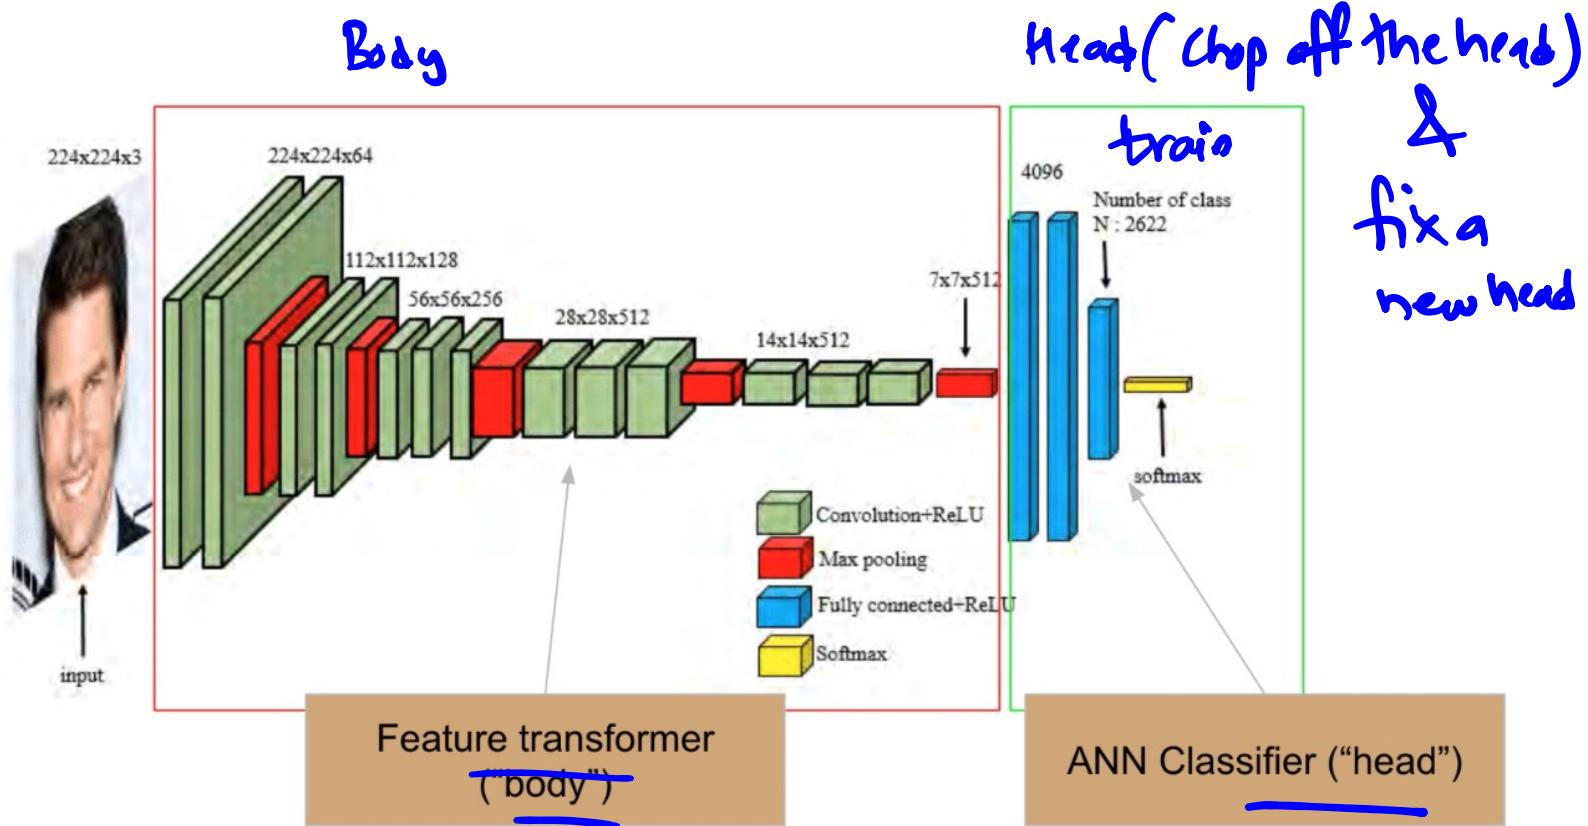


- Main idea: the features I found from one task may be useful for another task
- Transfer Learning took off in the field of computer vision
- ImageNet - Large-scale image dataset (millions of images, 1k categories)
- Because the dataset is so diverse, weights trained on this dataset can be applied to a large number of vision tasks
 - Cats vs Dogs
 - Cars vs Trucks
 - Even microscope images / images never seen before!



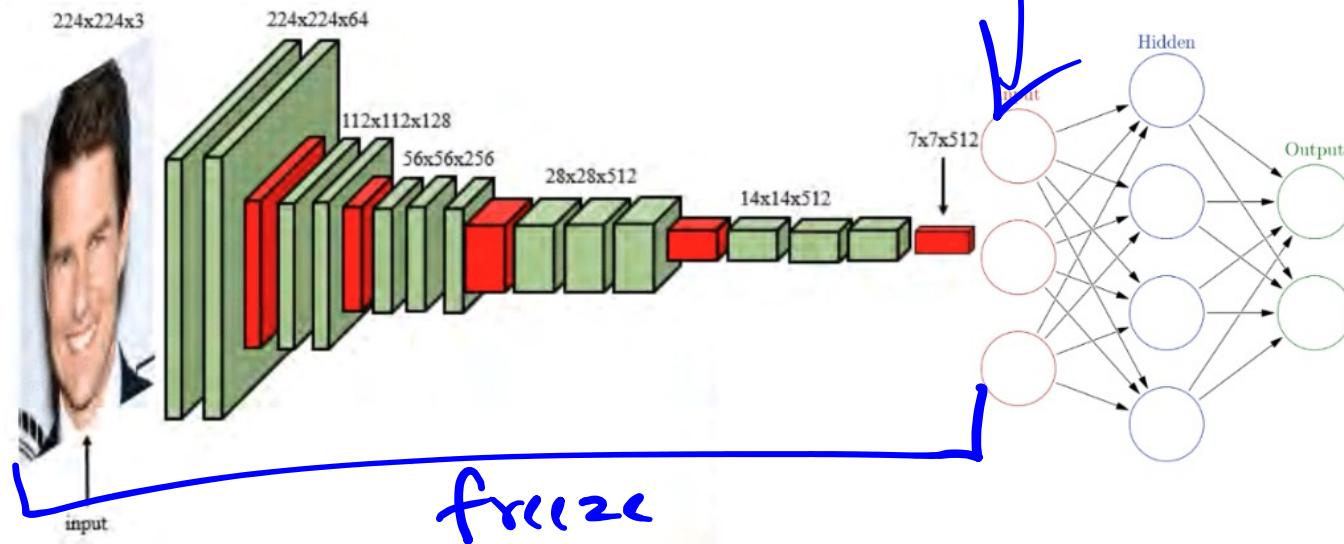
- It's not feasible for us to train on ImageNet ourselves
- Old days: training used to take days, weeks, even months!
- Now we can do it in minutes, using multi-GPU clusters, etc.
- The major CNNs which have won past ImageNet contests have publicly-released pre-trained weights
- And they are already included in Tensorflow/Keras!





Chop off the old "head", add a new head!
Can be logistic regression or ANN

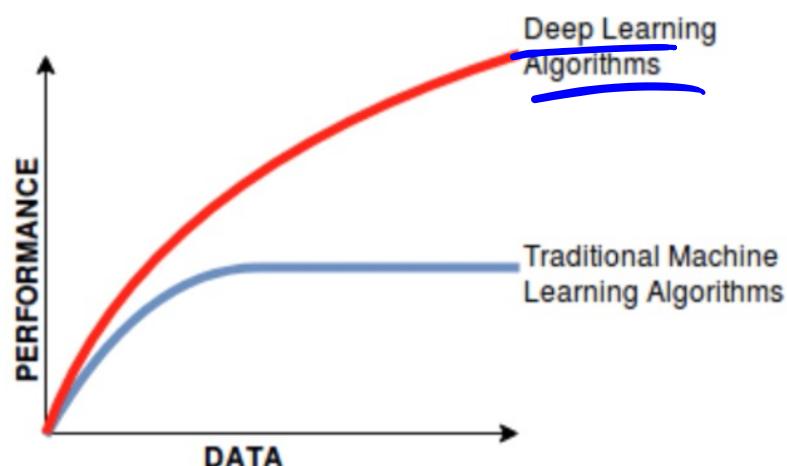
E.g. "cars vs trucks" may use a logistic regression with sigmoid



Don't a lot of data to build a state-of-the-art model

With transfer learning, this work has been done for us - the earlier features were already trained on lots of data (ImageNet)

Small dataset + a lot less weights helps us train **fast**



VGG

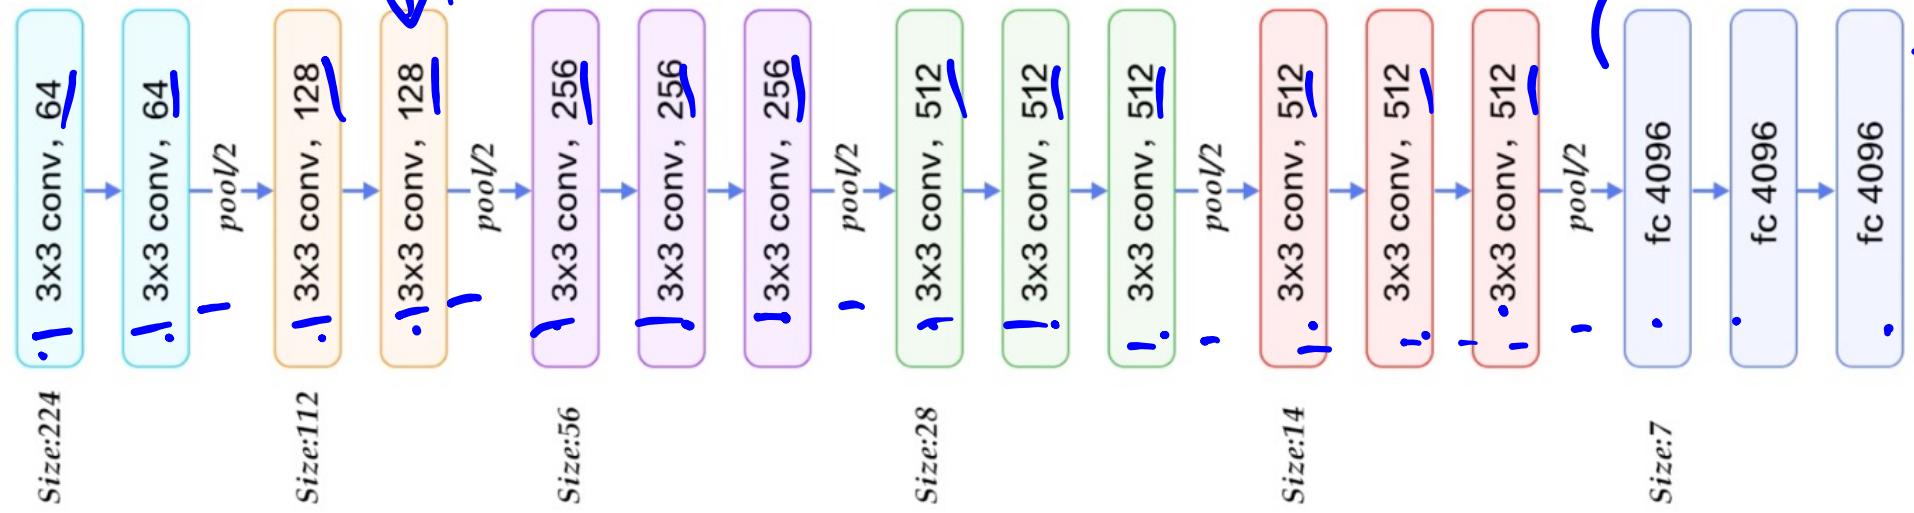
Named after the research group that created it - Visual Geometry Group

Not that different from CNNs we already know, just bigger

Options: VGG16, VGG19

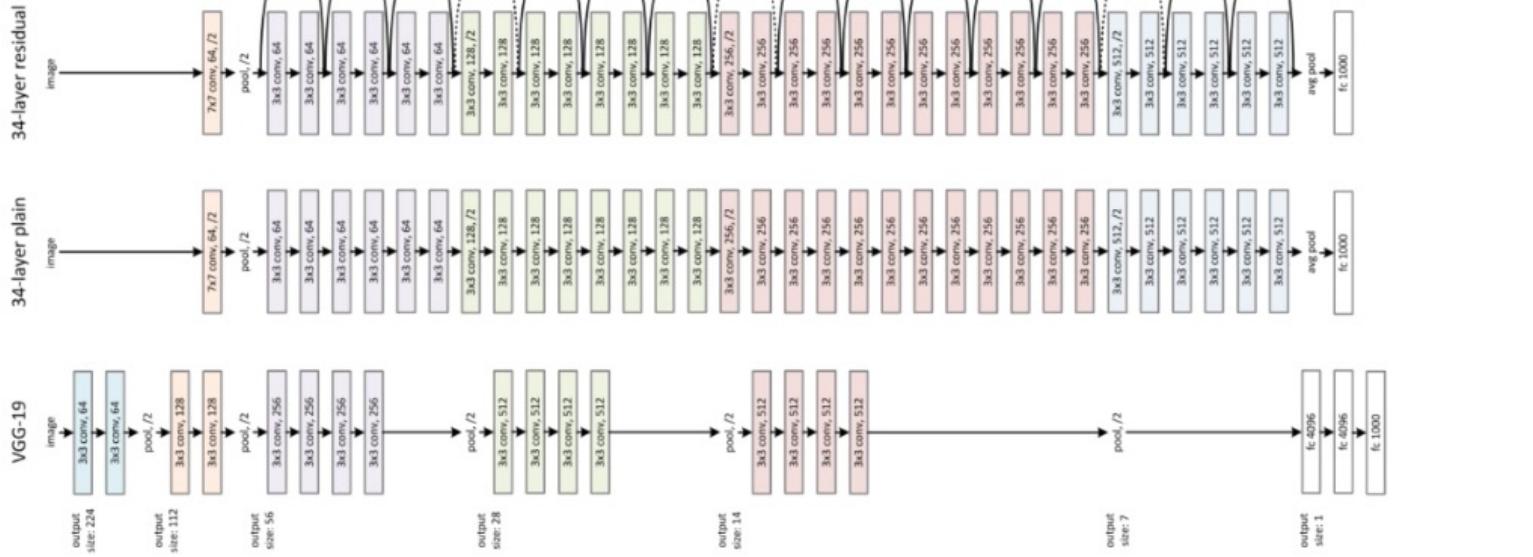
VGG

ANN



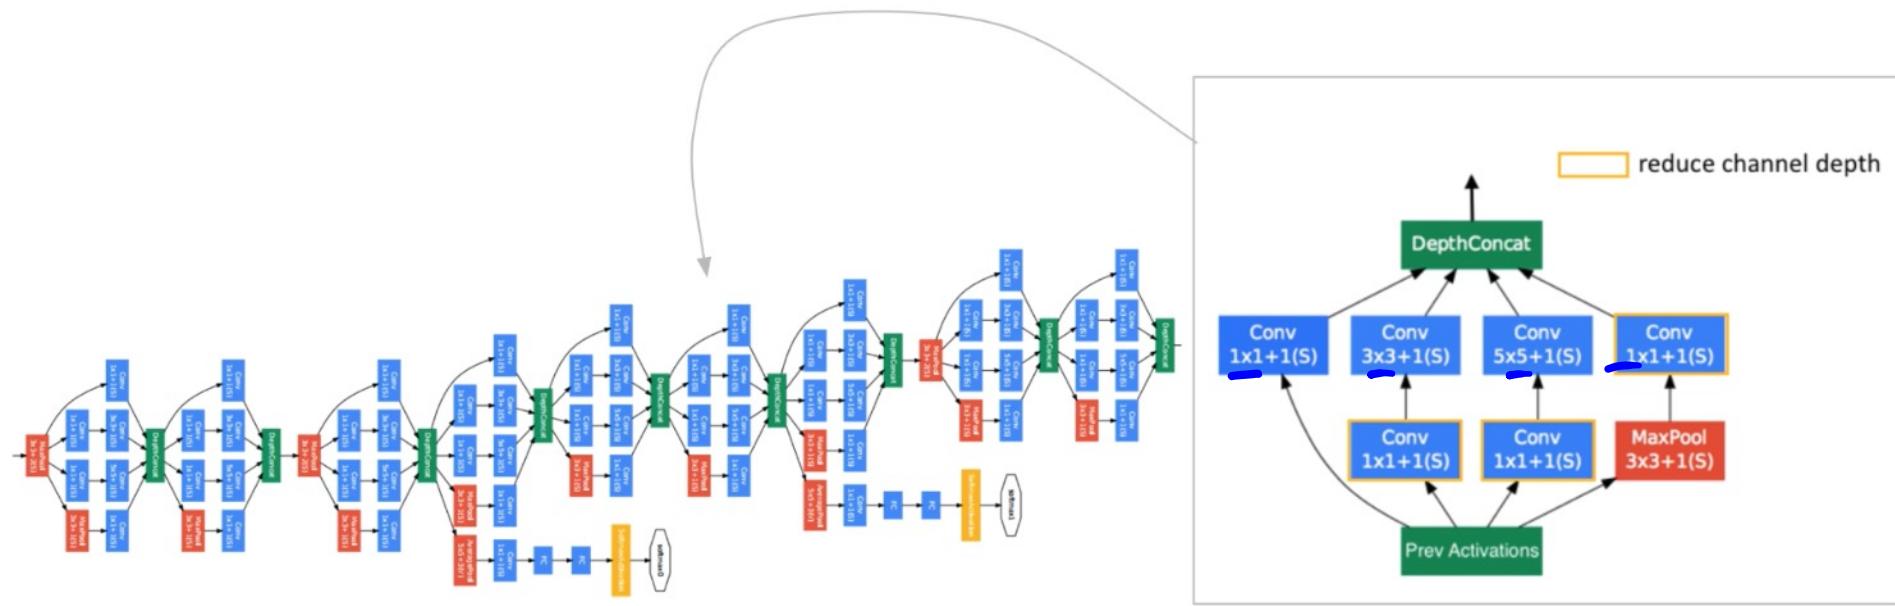
A CNN with branches (one branch is the identity function, so the other learns the *residual*)

Variations: ResNet50, ResNet101, ResNet152, ResNet_v2, ResNeXt



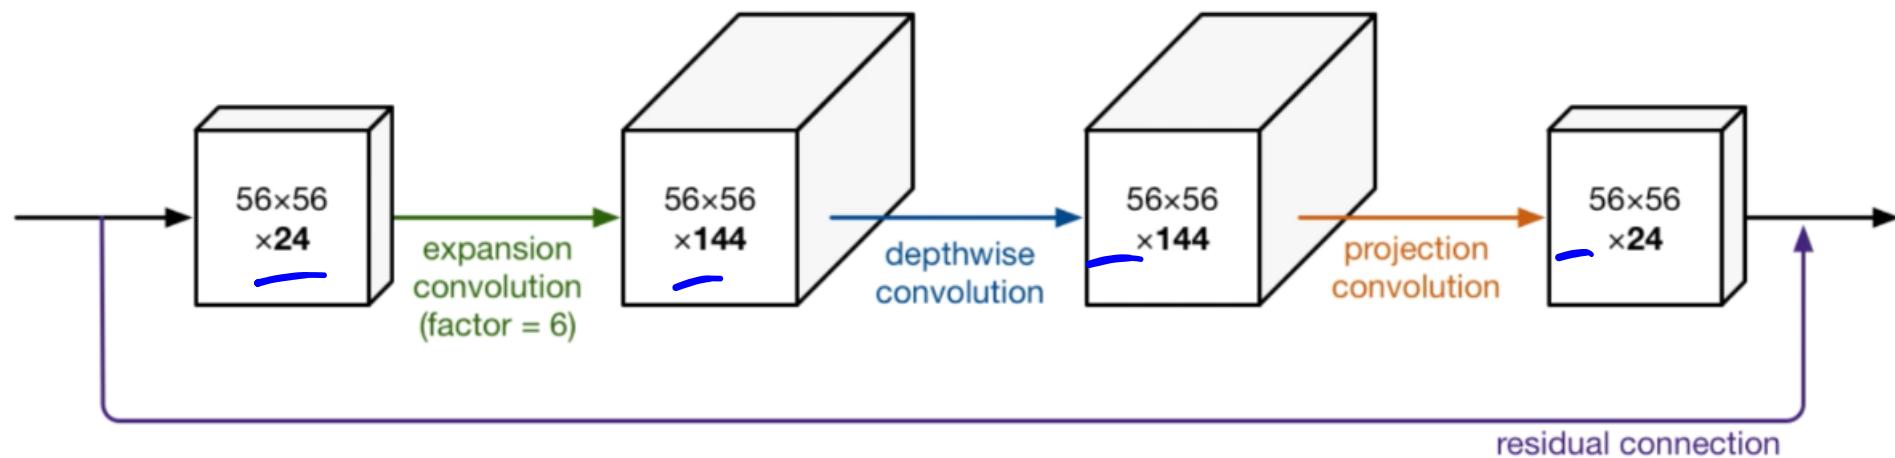
Multiple convolutions in parallel branches

Instead of trying to choose different filter sizes (1x1, 3x3, 5x5, etc.) just try them all!



MobileNet

Lightweight: makes a tradeoff between speed and accuracy
Meant for less powerful machines (mobile, embedded)



Since we are using pre-built CNNs, our data must be formatted just like the original

We usually work with RGB pixel values in [0, 1] or [-1, +1]

VGG (for example) uses BGR with pixel values centered but not scaled

Just import the preprocess_input function from the same module as your pretrained network

```
from keras.layers import Input, Lambda, Dense, Flatten  
from keras.models import Model  
from keras.applications.resnet50 import ResNet50, preprocess_input  
# from keras.applications.inception_v3 import InceptionV3, preprocess_input  
from keras.preprocessing import image  
from keras.preprocessing.image import ImageDataGenerator
```

When you are first learning about deep learning, it is convenient to have MNIST, CIFAR-10, SVHN, ... as a CSV / Numpy array

In the real world, images are not CSVs, images are images!

i.e. image **files** (JPEG, PNG, ...)

Real images are much larger than MNIST (28x28) and CIFAR-10 (32x32)

VGG and ResNet are trained on ImageNet images resized to 224x224



ImageNet

How much space does it take to store 1 million images of size 224 x 224?

1 million images x (224 x 224 x 3) bytes / image

~150 billion bytes

~140 GB

Would not fit in RAM on a standard machine



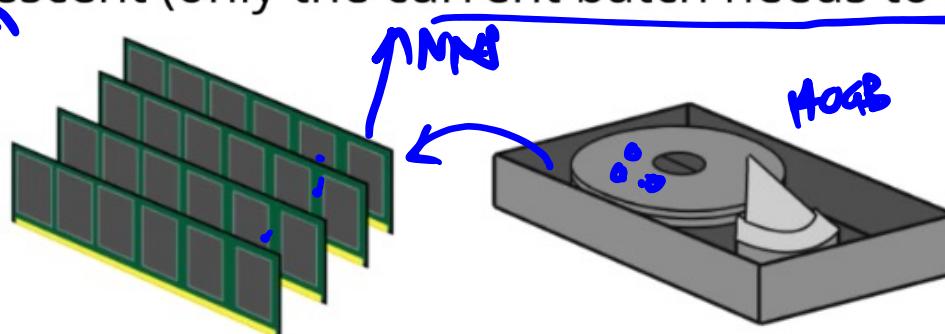
Approach the problem like an engineer, not that of a simple API user copying code off the Internet

Recognize the difference between disk and memory

Disk: slow, memory: fast

Model reads data from memory; image files live on disk

We do batch gradient descent (only the current batch needs to exist in memory)



MEMORY vs **STORAGE**

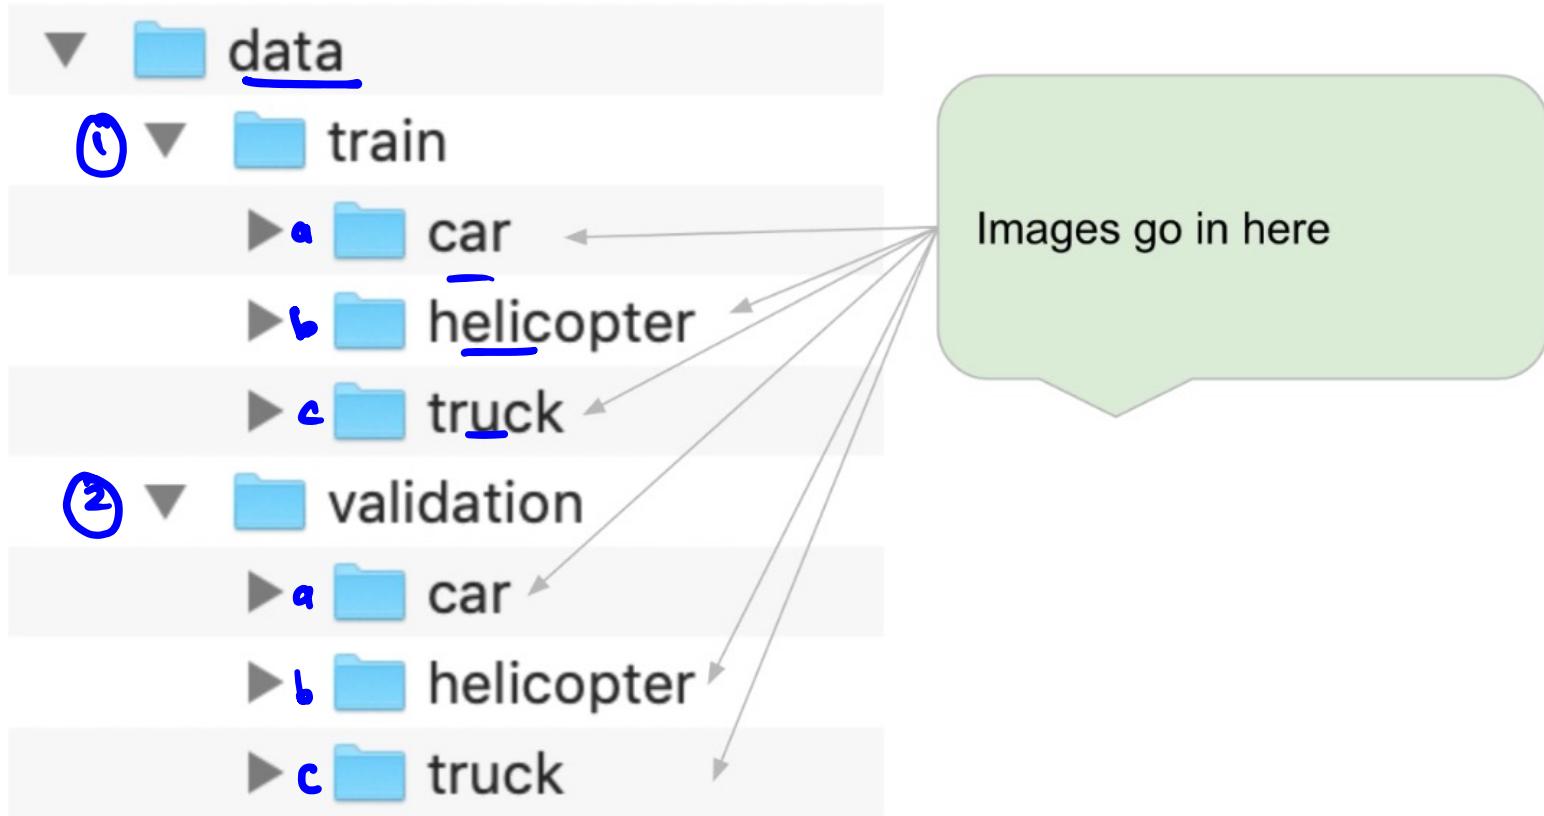
- Assume we have 2 arrays/lists: list of filenames, list of labels
- Suppose batch_size = 32

32000

n_batches = #BatchSize = 1000

for i in range(n_batches):
 x = load_images(filenames[i: i + 32])
 y = labels[i: i + 32]
 model.train_on_batch(x, y)

- The main ingredients:
- 1 `gen = ImageDataGenerator()`
 - Automatically generates data in batches
 - Data augmentation (shifting, rotation, flipping, ...)
 - Preprocessing (via preprocess_input)
- 2 `generator = gen.flow_from_directory()`
 - Specify target image size (e.g. resize all images to be 224 x 224), batch size
- 3 `model.fit_generator(generator)`
 - Used in place of model.fit(X, Y)



Approach #1 - use data augmentation with ImageDataagenerator

- Entire CNN computation must be inside loop

Approach #2 - precompute Z without data augmentation

- Only need to train a logistic regression on (Z, Y)



	With data augmentation <u> </u>	Without data augmentation <u> </u>
Speed	Slow (must pass through entire CNN)	Fast (only need to pass through 1 dense layer)
Generalization / Accuracy	Possibly better for generalization	Possibly worse for generalization

Reading the data
from disk over and
over is also slow!

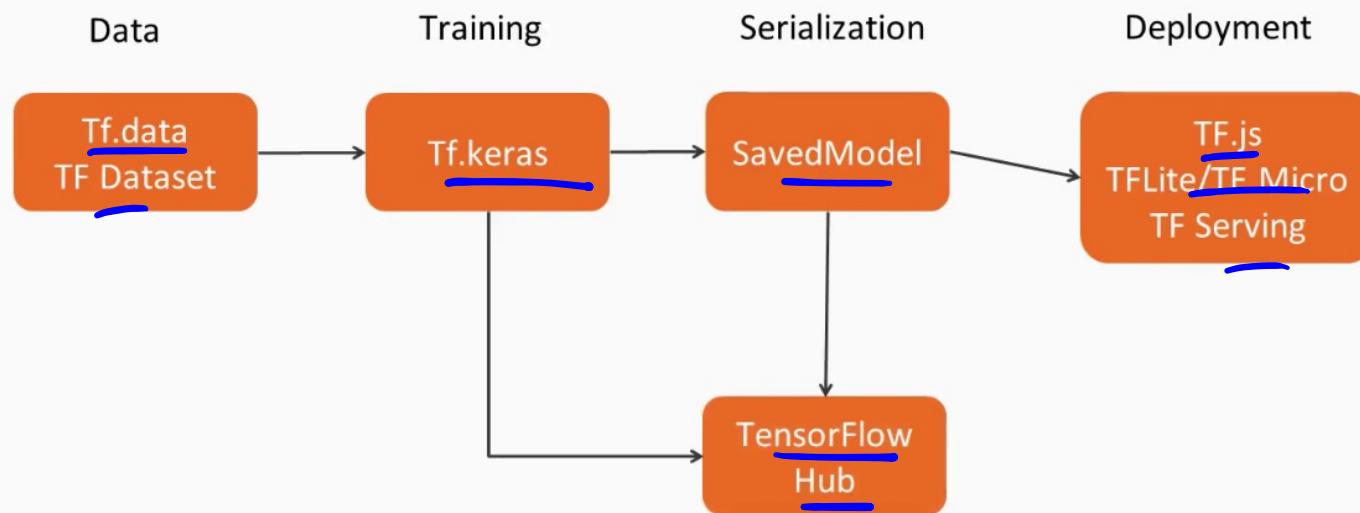
	With data augmentation	Without data augmentation
Speed	Slow (must pass through entire CNN)	Fast (only need to pass through 1 dense layer)
Generalization / Accuracy	Possibly better for generalization	Possibly worse for generalization

What Is TensorFlow Hub?

An open-source repository of
ready-to-use pre-trained models

- Announced at TF Dev Summit 2018
- Use the models directly

- More than just SavedModels on TensorFlow Hub
- TensorFlow.js
- TensorFlow Lite + Metadata
- Coral Edge TPU



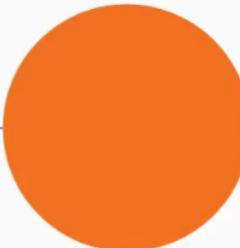
Image

- Classification
- Object detection
- Style Transfer
- Generator, and
so on



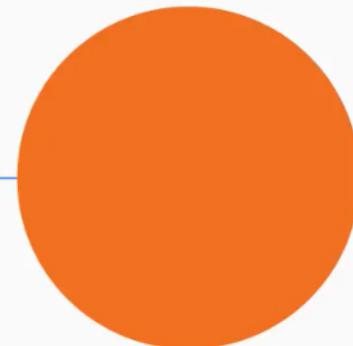
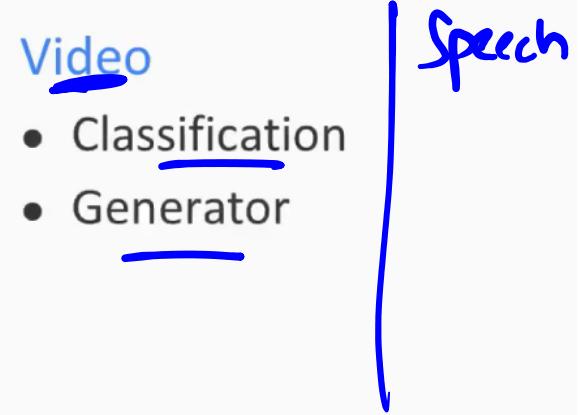
Text

- Classification
- BERT
- Embedding



Video

- Classification
- Generator

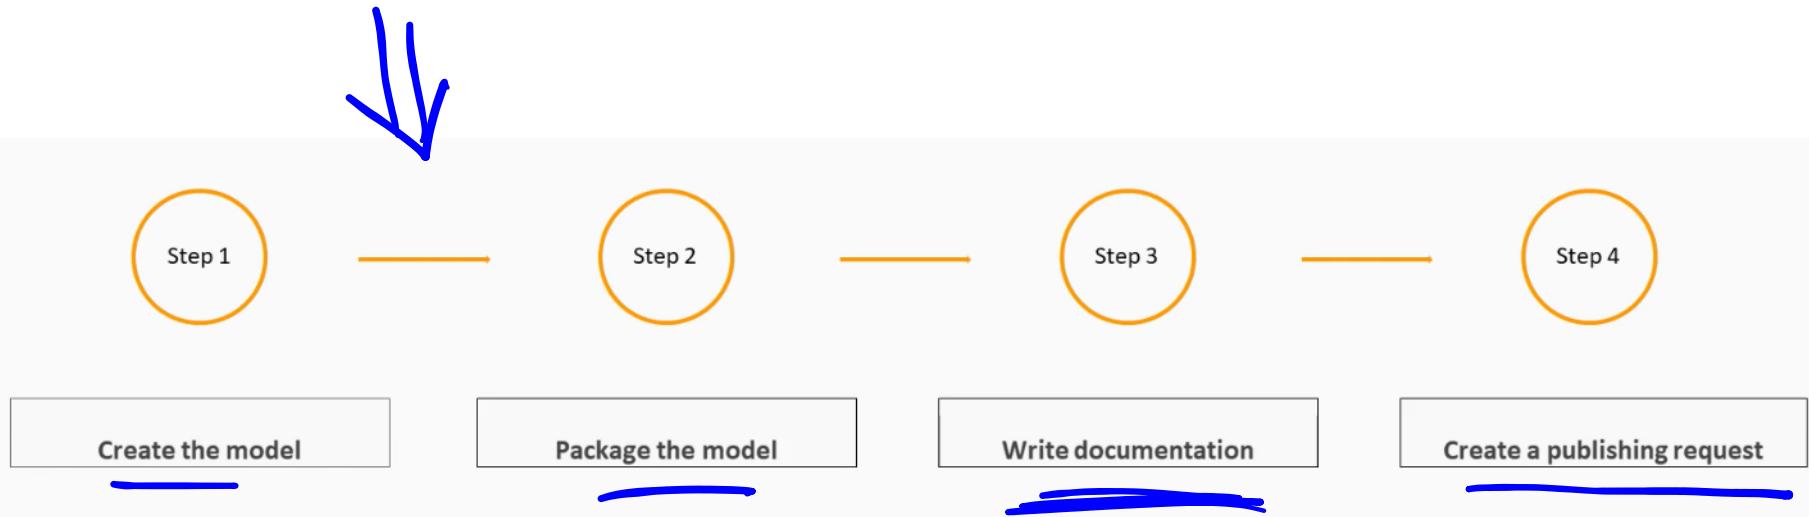




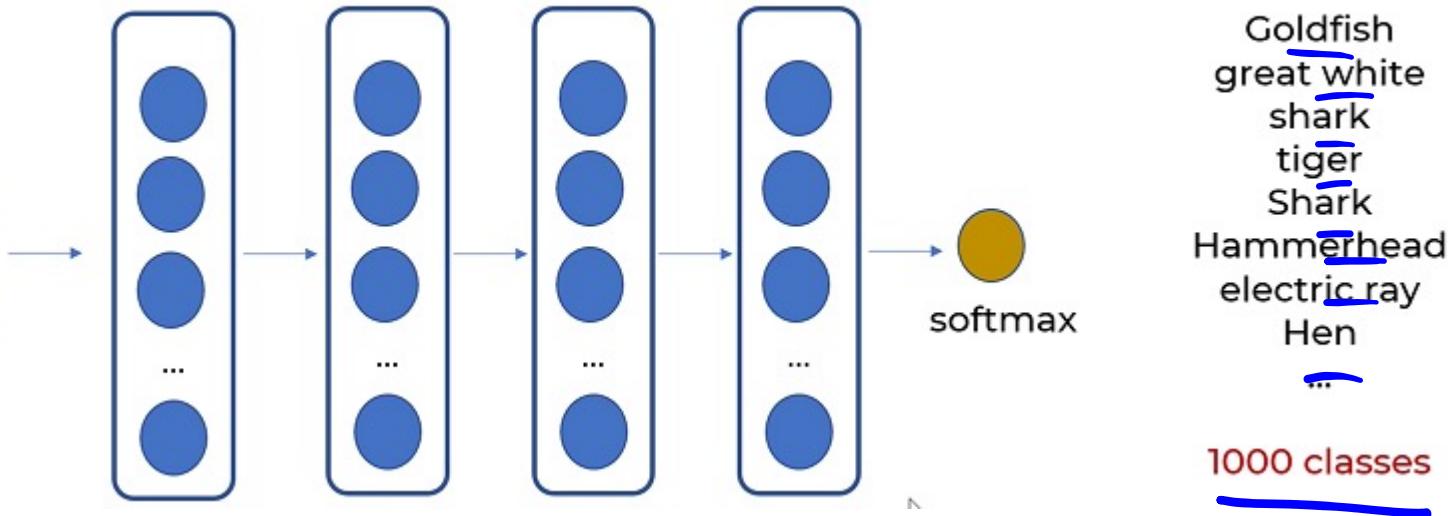
> 1000

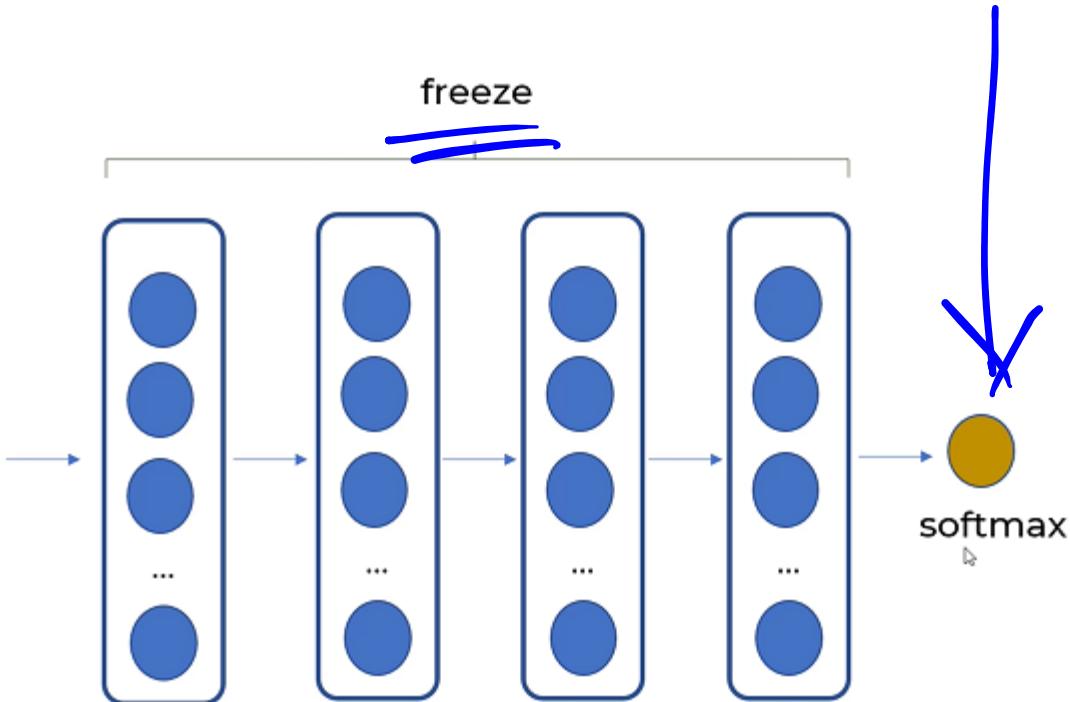
State of art models by Google and community

- Google
- DeepMind
- MediaPipe
- Kaggle
- NVIDIA
- Microsoft AI for earth
- The metropolitan Museum of Art
- iNaturalist



https://github.com/tensorflow/hub/tree/master/tfhub_dev





Flowers

Daisy
Rose
Tulip
Dandelion
sunflower

5 classes



#LifeKoKaroLift

Thank You!

Connect me: <https://www.linkedin.com/in/dr-darshan-ingle-corporate-trainer/>

